

Week 9: We can optimize that!

Amath 301

TA Session

Today

- 0. Cookbook optimization
- 1. Optimization suites in Matlab and Python
- 2. Techniques for ugly functions

UW is very strong in optimization/applied optimization:

Math 407/408/409, Math/Amath 514/515/516/518

EE 445/447/550/556/578, CSE 421/521/535/541

What's an optimization problem?

All optimization problems look like this:

$$\min_{x \in X} f(x)$$

Sometimes we can get $x^* = \arg \max_{x \in X} f(x)$ if we know x^* is unique.

Two pieces to work out:

1. A domain X of states to search over. Can be discrete, an interval, the real line, or much more complicated. Sometimes called the *Feasible Region*.
2. A function f called the "objective function. It's what we want to optimize.

Issues?

The optimal solution x^* may not exist or may not be unique. Mathematical analysis is needed to **prove** a unique optimal solution exists.

Sometimes the work is done for us already

Common domains: $[0, 1]$, $[-1, 1]$, \mathbb{R} , \mathbb{R}^n , \mathbb{Z} (harder) but not always.

Some useful facts/results:

1. If f is continuous and X is a closed interval, then at least one optimal solution exists.
2. If f is convex over X and X is a closed interval or \mathbb{R} , then a **unique** optimal value exists.

For twice-differentiable functions, it suffices to show that $f''(x) > 0$ for all $x \in X$ to claim convexity. To prove properties using convexity, we often restrict the domain.

The reality

Many optimization problems can be solved over an entire domain. Sometimes, you need more structure to your solution.

Constraints: Additional equations you require your solution to satisfy. These come from your problem domain (e.g. conservation of mass, energy, etc).

Examples:

- $x_1 + x_2 + x_3 = 1$ Conservation laws can be enforced
- $x_1 + ex_2 - x_3^2 = -\pi$ Do not need to be "nice"
- $x_1 > 0$ Positivity or non-negativity can be enforced
- $0 \leq x \leq 1$ (x lies in $[0, 1]$)

Why just min? Why not max too?

$$\min_{x \in X} f(x) \quad \text{vs} \quad \max_{x \in X} f(x)$$

Optimization and curve fitting?

Fitting a curve is as easy (or as hard) as solving an optimization problem.

- Curve fitting - minimize an error term.
 - you choose (or are told) the form of the error term
 - You choose (or are told) the form of the curve

Example:

$$E(a) = \sqrt{\frac{1}{n} \sum_{j=1}^n (f(x_j; a) - y_j)^2}$$

where f is the model you want to fit.

Let's write some code.

In code:

```
def sample_model(x, params):  
    """ A vectorized model"""  
    return params[0] * np.sin(params[1] * x)  
  
def E2error(x, xdata, ydata, model):  
    """ Given parameters `x`, some data, and a model, return the l2 error"""  
    n = xdata.size # how many points do you have? - 1d only  
  
    # Using your parameters and a model, compute the predictions  
    preds = model(xdata, x)  
    error = np.sqrt(np.sum((preds - ydata) ** 2)/n)  
    return error  
  
init_guess = [1,1]  
opt_params = fmin(lambda x: E2error(x, xdata, ydata, sample_model), init_guess)
```

In code:

```
function preds = simple_model(x, params):  
    preds = params(1) * np.sin(params(2)) .* x);  
end  
  
function error = E2error(x,xdata, ydata, model):  
    % Given parameters `x`, some data, and a model, return the l2 error  
    n = length(xdata); % how many points do you have? - 1d only  
  
    # Using your parameters and a model, compute the predictions  
    preds = model(xdata, x); % call the model  
    error = sqrt(sum((preds - ydata) .^ 2)./n); % error calculation  
end
```

```
init_guess = [1,1];  
opt_params = fminsearch(@(x) (@E2error(x, xdata, ydata, @simple_model)), init_guess)
```

Optimization tools

Use `fmin` or `fminsearch` for HW7.

"Traditional" optimization solvers require 2,3, or 4 pieces of information usually:

1. Initial guess
2. Function f to optimize
3. Region to search over
4. Derivative information (Jacobian or Hessian matrices) - matrices of derivatives for multi-valued functions

Always need 1. and 2., sometimes 3., and sometimes 4.

Sometimes the initial guess is easy - ODE/PDE problems, or 0. Others (ML, data sci), not so easy!

Python

The more complicated your method or problem, the more code you write.

- `Scipy.optimize` module (Start here for your problems). Modern API for new code
 - (local) `minimize` function offers 14 different algorithms
 - (local) `minimize_scalar` offers 3 methods for single variable problems.
 - (global) `shgo`, `basinhopping`, `brute` (force!), `dual_annealing`

Matlab

- (local) Optimization toolbox:
 - `fminsearch` Unconstrained problems (5 algorithm choices)
 - `fmincon` Constrained problems (5 algorithm choices)
- Global optimization toolbox:
 - `particleswarm` [Video](#)
 - `simulannealbnd` Simulated annealing [Video](#), [Demo](#)

What to do if these techniques don't work?

Two options:

1. Deep learning. Add a few more million parameters and you can fit just about any data.
2. Monte Carlo Methods. A broad class of stochastic bootstrapping on statistical methods that leverage the randomness to find meaningful patterns. The computational sibling to computational Bayesian statistics.
 - Major application: Generate samples from "interesting" distributions (data-driven).

To speed these up, get a larger (super) computer.

Deep learning and Machine learning

Python

- `sklearn` (`conda install scikit-learn -y`) A more machine learning focused library
- `TensorFlow` and `PyTorch` - Deep learning libraries. Now feature numpy-like APIs
 - Interfaces to C++ and CUDA code that handle the actual computation

Matlab

- Statistics and Machine Learning Toolbox
- Deep learning toolbox

How does deep learning work?

Gradients. The backprop algorithm is just fancy chain rule.

These calculations are handled for you by Tensorflow/PyTorch/Matlab.

